

VARC Hint Guidance for CFF2

Skef Iterum

March 4, 2024

CFSH — Compact Font Format Supplementary Hint Table

(This section is added as a new table specification.)

‘CFSH’ is an optional table used to provide supplementary ‘CFF2’-format PrivateDICT structures for VARC composite glyphs, and to map each VARC glyph to those PrivateDICTs or to PrivateDICTs in the ‘CFF2’ table. Its contents are certain CFF2 subtables, in some cases slightly modified.

CFSH Header

Type	Name	Description
uint16	major_version	Table major version number (=1)
uint16	minor_version	Table minor version number (=0)
Offset32	privateDICTIndexOffset	Offset (from start of CFSH table) to a ‘CFF2’ INDEX of Private DICTs. 0 if no PrivateDICTIndex.
uint16	initialPrivateDICT	The FontDICT index associated with of the first entry in the PrivateDICTIndex (default is 0).
Offset32	fdSelectOffset	Offset (from start of CFSH table) to the FontDICTSelect subtable. Must not be 0.
Offset32	itemVarStoreOffset	Offset (from start of CFSH table) to the Item Variation Store table (may be 0)

Private DICT Index and initialPrivateDICT

In the CFF2 table a Font DICT INDEX contains FontDICT structures, each of which encodes the size and offset of a Private DICT as is described in (crossreference section on CFF2 Private DICTs). CFSH also encodes CFF2 Private DICTs but eliminates the indirection through the FontDICT. Instead, it contains a Private DICT Index, which is a CFF2 Index structure (cross-reference section on CFF2 Indexes) that stores a list of Private DICTs directly.

The initialPrivateDICT field is the “FontDICT index” associated with the first entry in the Private DICT Index. Thus if initialPrivateDICT is 24, the index of the first CFSH private dict is 24, the index of the second is 25, and so on. The initialPrivateDICT field is normally be set to the size of the CFF2 FontDICT INDEX, so that the index of the first CFSH Private DICT is one greater than the index of the last CFF2 Private DICT. The indexes of the two sets of Private DICTs must not overlap.

If the itemVarStoreOffset field is non-zero, then the vsindex and blend operators relate to the Item Variation Store it points to. If the field is zero then those operators relate to the Item Variation Store in the CFF2 table.

A PrivateDICT in the CFSH table must not include the Subrs operator.

The FontDICTSelect Offset

The FontDICTSelect offset points to a CFF2 FontDICTSelect subtable (cross-reference ‘CFF2’ section on FDSelect). A client should ignore this field when it is set to 0 to support future minor extensions of the table. However, in a version 1.0 ‘CFSH’ table the offset should not be 0. As of version 1.0 only FontDICTSelect format 4, as described in (cross-reference section on FDSelect format 4), is supported, with one modification: it is not required that the fd field in the first Range4 record be 0.

The FontDICTSelect subtable in ‘CFSH’ can overlap with the FontDICTSelect subtable in the ‘CFF2’ table but there must not be gap between the last glyph mapped in ‘CFF2’ and the first glyph mapped in ‘CFSH’. The sentinel field in ‘CFSH’ must be the highest GID defined in the font (from maxp or MAXP).

The itemVarStore Offset

When not zero, this field points to an Item Variation Store used for the vsindex and blend operators for any PrivateDICTs in the CFSH table.

VARC and Hinting

(This will be a new, appropriately placed section in the VARC chapter.)

When a VARC composite glyph is built from ‘glyf’ components that include TT instructions, or from ‘CFF2’ components that include hinting parameters, that data should be considered part of the composite. Whether and how hint data is used when rasterizing the glyph can depend on a number of factors, including the transforms applied to the component both within VARC or “externally” (e.g. using a CSS transform).

The ‘glyf’ table has long included its own composite format that makes use of the TT instructions of component glyphs. That mechanism can serve as a model for when and how to apply instructions when rasterizing a ‘glyf’-based VARC composite.

Neither CFF nor CFF2, in contrast, has internal support for compositing, and while the ‘COLR’ table implicitly adds such support, COLR leaves the question of hinted rasterization open. The rest of this section closes this gap by clarifying how to adapt ‘CFF2’ hinting parameters when rasterizing a VARC composite glyph. Those parameter can then be applied as they would be when rasterizing a hinted CharString in a CFF2 table.

Hinting CFF2 components in a VARC context

The hinting parameters in the ‘CFF2’ table evolved from related information included in PostScript Type 1 fonts, and consists of a combination of PrivateDict parameters—each of which is associated with a particular subset of glyphs—and per-glyph parameters. With VARC the PrivateDict parameters are taken from the PrivateDICT associated with the composite glyph—or if there are multiple layers of compositing, the PrivateDICT associated with the outer-most composite. The per-glyph parameters are adapted from the hint operators included in the CFF2 CharString of a

component glyph, and therefore in VARC are technically per-component rather than per-glyph. These must be adapted to account for the transformations and translations applied in each layer of compositing.

The Private DICT

When a CFF2 glyph is included in a VARC composite as a component, both the PrivateDICT of the component and that of the composite must typically be consulted in order to render it.

The PrivateDICT of the component glyph will always be in the CFF2 table and is mapped via the component's GID in the CFF2 FontDICTSelect subtable (unless all CFF2 glyphs use the same PrivateDICT, in which case there will only be one). That PrivateDICT may contain Subrs or vsindex operators needed to desubroutinize the component's CharString and to resolve any blends it contains relative to the specified location in design space.

The PrivateDICT of the (top-level) VARC composite glyph will either be in the CFF2 table or the CFSH table. It is found using this procedure:

1. If there is a CFSH table and it contains a FontDICTSelect structure, that is checked for the GID of the VARC composite glyph. If the GID is mapped, the PrivateDICT index it maps to is the index for the composite.
2. If there is no CFSH table, no FontDICTSelect structure in the table, or the GID of the composite is not mapped in that structure, the FontDICTSelect structure in the CFF2 table is checked for the GID, and the PrivateDICT index it maps to is the index for the composite.
3. If neither FontDICTSelect structure maps the GID the font is malformed. If there is a CFSH table, the value of the initialPrivateDICT field is the fallback index for the composite. Otherwise the fallback index is 0.
4. If there is a CFSH table and the index for the composite is greater than or equal to its initialPrivateDICT field, the value of that field is subtracted from the index and the result used as the offset into the CFSH PrivateDICT Index.
5. Otherwise the index for the composite is used as the offset into the CFF2 FontDICT Index, and its PrivateDICT is used for the composite.

Per-glyph parameters

A hinted glyph has some combination of these parameters:

1. Horizontal and vertical stem regions (hstem(hm), vstem(hm))
2. Hintmasks
3. Counter hinting (cntrmask)

These are described in (cross-reference section(s) on per-glyph hinting parameters).

Relative to a given “top-level” composite glyph to be rendered, an “atomic” (or “bottom-level”) component glyph may be subject to multiple sets of transformations and translations, with one set per level of VARC compositing. How these transformations affect component point values is documented elsewhere. This section describes how to apply analogous transformations and translations to stem regions, and how to decide whether to proceed with hinting in a given dimension of the component glyph (horizontal or vertical) or to *cancel* hinting in that dimension.

Intuitively, hinting in a given dimension should be cancelled when there is rotation that is not a multiple of 180 degrees, or when there is any skew in the opposite dimension.

(Note that in addition to VARC composite transforms a glyph may also be subject to “external” transforms, such as those specified with a Cascading Style Sheets transform property. A rasterizer implementation can either avoid hinting in such circumstances or compose those transforms into the cumulative transformation of the stem hints as described below.)

Whether hinting in a given dimension should be canceled, and how the stems in that dimension should be adjusted when it is not cancelled, can be determined by performing the cumulative transformation on the three points p_1 (100, 0), p_2 (0, 0) and p_3 (0, 100) and considering the result. If we call the transformed points p_1' , p_2' , and p_3' respectively, adjustment of the horizontal stems proceeds as follows:

If the line from p_1' to p_2' is not close enough to horizontal for hinting purposes (which may vary by rasterizer implementation), hinting of horizontal stems is cancelled. Otherwise hinting proceeds in that dimension with a scale factor of

$$s = (y_3' - y_2') / 100$$

and a translation factor of

$$t = y_2'$$

If hinting proceeds, the horizontal stem deltas are unpacked into positional bottom, top pairs and each pair is processed as follows:

1. If the pair does not represent an edge stem, each value is multiplied by s and translated by t . If s is negative the top and bottom values are swapped.
2. If the pair represents a bottom edge stem, the lower edge is multiplied by s and translated by t . If s is positive the pair is re-encoded as a bottom edge stem with the adjusted value as the bottom. If s is negative the pair is re-encoded as a top edge stem with the adjusted value as the top.

For example, if the original bottom position was b , the unpacked pair was encoded as having a width of -21 (i.e. the first number unpacked to $b + 21$ (the “upper edge”) and the second to b (the “lower edge”). If s is positive the adjusted lower edge will be $b*s + t$ and the adjusted upper edge will be $b*s + t + 21$. (If necessary this can be re-encoded as a bottom edge hint with the first value $b*s + t + 21$ and a width of -21.)

If s is negative then the hint is adjusted to be a top hint, with $b*s + t$ as the upper edge and $b*s + t - 20$ as the lower edge, which can be re-encoded as an initial value of $b*s + t$ and a width of -20.

3. If the pair represents a top edge stem, the upper edge is multiplied by s and translated by t . If s is positive the pair is re-encoded as a top edge stem with the adjusted value as the top. If s is negative the pair is re-encoded as a bottom edge stem with the adjusted value as the bottom.

If s is positive the adjusted stem pairs simply replace the respective original pairs. If s is negative the orders of both the stem pairs and the corresponding bits in each of the hintmasks and cntrmasks for the glyph are reversed, and those stem pairs and masks are used to render the component.

Adjustment or cancellation of vertical stems is analogous to the horizontal case.

Adjustments to CFF2 chapter

(These aren't quite in "editorial form" yet, because the previous proposal that extensively modified the CFF2 text is not yet integrated.)

1. When a font has a VARC table, the length of the CFF2 CharStringINDEX can be less than the maxp count, otherwise it must be equal to the maxp count.
2. When a font has a VARC table, the highest glyph mapped by the FontDICTSelect structure can be less than the maxp count as long as two requirements are met. The first requirement is that there is a CFSH table that maps any remaining glyphs. The second requirement is that every glyph in the CharStringINDEX must be mapped.

The ranges of mapped glyphs can overlap as described in (cross-reference CFSH section "The FontDICTSelect Offset")